

---

# Internetworking with Sockets

# Cross-host Interprocess Communication (IPC)

- ▶ Typically client-server model over network
- ▶ Server - Provides a service
- ▶ Server - Waits for clients to connect
- ▶ Clients - Connect to utilize the service
- ▶ Clients - Possibly more than one at a time

# The Internet Protocol

- ▶ Each device in a network is assigned an IP address
- ▶ IPv4 32 bit, IPv6 128 bit
  - IPv4 (in dec)  
69.89.31.226  $\Leftarrow$  4 octets
  - IPv6 (in hex)  
2001:0db8:0a0b:12f0:0000:0000:0000:0001  $\Leftarrow$  8 16-bit blocks
- ▶ Each device may host many services
- ▶ Accessing a service requires a (IP,port) pair
- ▶ Services you know of: ssh (port 22), http (port 80), DNS (port 53), DHCP (ports 67,68)

# Common Service Use Cases

## Browse the World Wide Web

- ▶ Each device has a static IP
- ▶ DNS used to translate `www.google.com` to `216.58.213.4`
- ▶ Contact service at `216.58.213.4` and port 80 (`http`)

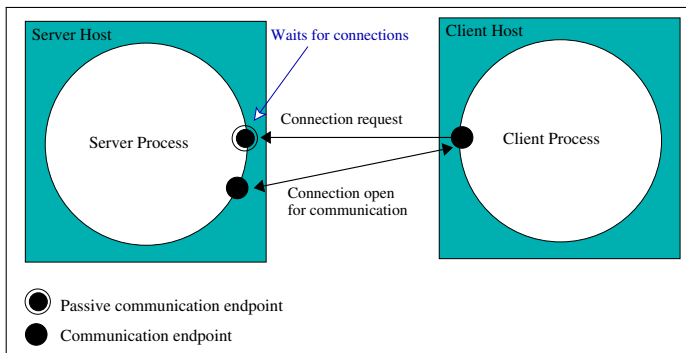
## Common Service Use Cases

Your home network.

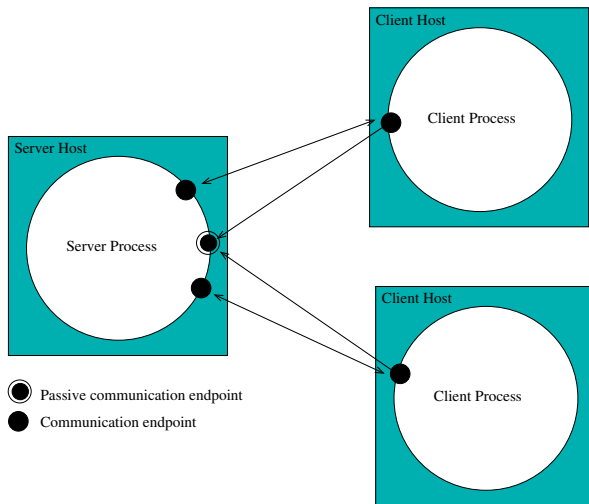
- ▶ You turn on your modem. It gets a public from you ISP (eg. 79.166.80.131)
- ▶ Your modem runs a DHCP server giving IPs in 192.168.x.y
- ▶ Your modem acts as a Internet gateway. Translates IPs from 192.168.x.y to 79.166.80.131. IP Masquerade.
- ▶ What if you need to setup a service running inside your 192.168.x.y network available to the internet?  
Do port forwarding.

# The Transmission Control Protocol

- ▶ TCP Uses acknowledgments
- ▶ Non-acknowledged messages are retransmitted
- ▶ Messages re-ordered by the receiver's OS network stack
- ▶ Application sees a properly ordered *data stream*



## TCP - multiple clients



# Sockets

- ▶ A socket is a communication endpoint
- ▶ Processes refer to a socket using an *integer descriptor*
- ▶ Communication domain
  - ▶ Internet domain (over internet)
  - ▶ Unix domain (same host)
- ▶ Communication type
  - ▶ Stream (usually TCP)
  - ▶ Datagram (usually UDP)



## TCP vs. UDP

	TCP	UDP
Connection Required	✓	✗
Reliability	✓	✗
Message Boundaries	✗	✓
In-Order Data Delivery	✓	✗
Socket Type	SOCK_STREAM	SOCK_DGRAM
Socket Domain	Internet	Internet
Latency	higher	lower
Flow Control	✓	✗

## Serial Server (TCP)

Create listening socket *a*

**loop**

Wait for client request on *a*

Open two-way channel *b* with client

**while** request received through *b* **do**

Process request

Send response through *b*

**end while**

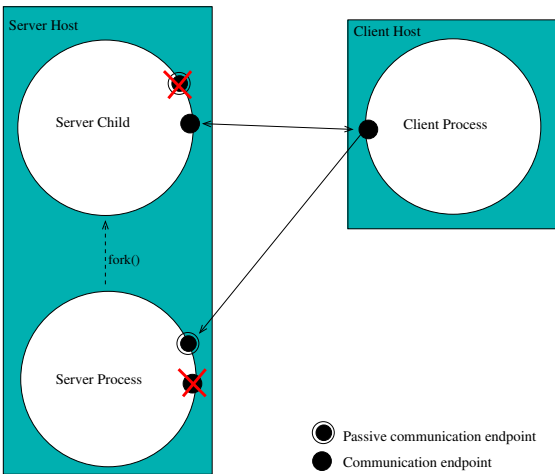
Close file descriptor of *b*

**end loop**

Drawbacks:

- ▶ Serves only one client at a time
- ▶ Other clients are forced to wait or even fail

# 1 process per client model



- ▶ New process forked for each client
- ▶ Multiple clients served at the same time
- ▶ Inefficient, too many clients → too many processes

# 1 process per client model

## *Parent process*

Create listening socket  $a$

### **loop**

Wait for client request on  $a$

Create two-way channel  $b$  with client

Fork a child to handle the client

Close file descriptor of  $b$

### **end loop**

## *Child process*

Close listening socket  $a$

Serve client requests through  $b$

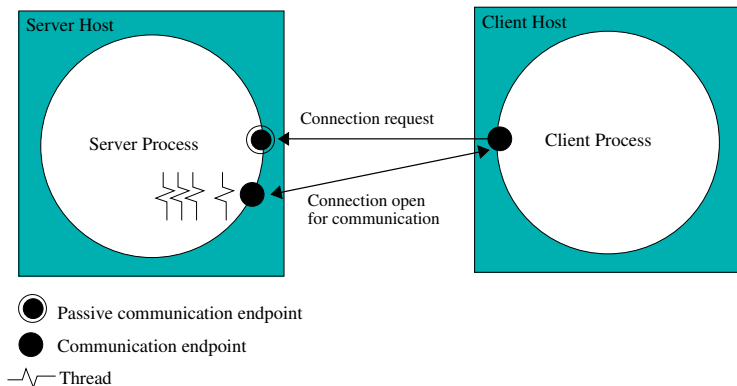
Close private channel  $b$

Exit

## Parent process: why close file descriptor *b*?

- ▶ Parent doesn't need this file descriptor
- ▶ Risk of running out of file descriptors otherwise
- ▶ Enables the destruction of the channel once the other two parties (child & client) close their file descriptors
- ▶ Enables the child process to receive EOF after the client closes its end of the channel (and vice versa).

## Multithreaded server model



- ▶ Multiple threads handle multiple clients concurrently
- ▶ Drawback: Requires synchronization for access to shared resources

## Dealing with byte order

- ▶ Byte order poses a problem for the communication among different architectures.
- ▶ Network Protocols specify a byte ordering: ip addresses, port numbers etc. are all in what is known as *Network Byte Order*
- ▶ Convert long/short integers between *Host* and *Network Byte Order*

```
/* host to network byte order for long -32bits */
uint32_t htonl(uint32_t hostlong);
/* host to network byte order for short -16bits */
uint16_t htons(uint16_t hostshort);
/* network to host byte order for long -32bits */
uint32_t ntohl(uint32_t netlong);
/* network to host byte order for short -16bits */
uint16_t ntohs(uint16_t netshort);
```

## Depicting the Byte Order ByteOrder-p16.c

```
#include <stdio.h>
#include <arpa/inet.h>

int main(){
    uint16_t nhost = 0xD04C, nnetwork;
    unsigned char *p;
    p=(unsigned char *)&nhost;
    printf("%x %x \n", *p, *(p+1));
    /* 16-bit number from host to network byte order */
    nnetwork=htons(nhost);
    p=(unsigned char *)&nnetwork;
    printf("%x %x \n", *p, *(p+1));
    exit(1);
}
```

- Experimenting with an Intel-based (Little-Endian) machine:

```
ad@haiku:~/src$ ./ByteOrder-p16
4c d0
d0 4c
ad@haiku:~/src$
```

- Experimenting with a Sparc (Big-Endian/Network Byte Order) machine:

```
pubsrv1:/k24-examples> ./ByteOrder-p16
d0 4c
d0 4c
pubsrv1:/k24-examples>
```



## From *Domain Names* to *Addresses* and back

- ▶ An *address* is needed for network communication
- ▶ We often have to *resolve* the address from a domain name.  
ex. spiderman.di.uoa.gr  $\leftrightarrow$  195.134.66.107

```
struct hostent {  
    char    *h_name;           /* official name of host */  
    char    **h_aliases;       /* aliases (alt. names) */  
    int     h_addrtype;        /* usually AF_INET */  
    int     h_length;          /* bytelength of address */  
    char    **h_addr_list;     /* pointer to array of network addresses */  
};  
  
struct hostent *gethostbyname(const char *name);  
  
struct hostent *gethostbyaddr(const void *addr, socklen_t len, int type);
```

- ▶ For error reporting use `h_error` & `hstrerror(int err)`.
- ▶ Both calls return pointers to statically allocated `hostent` structure on success and `NULL` on error.

# Resolving names for machines

```
#include <netdb.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

void main(int argc, char **argv){
    int i=0;
    char hostname[50], symbolicip[50];
    struct hostent *mymachine;
    struct in_addr **addr_list;

    if (argc!=2 ) {printf("Usage: GetHostByName-p18 host-name\n"); exit(0);}

    if ( (mymachine=gethostbyname(argv[1])) == NULL)
        printf("Could not resolved Name: %s\n",argv[1]);
    else
    {
        printf("Name To Be Resolved: %s\n", mymachine->h_name);
        printf("Name Length in Bytes: %d\n", mymachine->h_length);
        addr_list = (struct in_addr **) mymachine->h_addr_list;
        for(i = 0; addr_list[i] != NULL; i++) {
            strcpy(symbolicip , inet_ntoa(*addr_list[i]) );
            printf("%s resolved to %s \n",mymachine->h_name,symbolicip);
        }
    }
}
```

# Resolving names

```
ad@haiku:~/src$  
ad@haiku:~/src$ ./GetHostByName-p18 federal.gov.ar  
Name To Be Resolved: federal.gov.ar  
Name Length in Bytes: 4  
federal.gov.ar resolved to 190.210.161.110  
ad@haiku:~/src$  
ad@haiku:~/src$  
ad@haiku:~/src$ ./GetHostByName-p18 www.bbc.co.uk  
Name To Be Resolved: www.bbc.net.uk  
Name Length in Bytes: 4  
www.bbc.net.uk resolved to 212.58.246.95  
www.bbc.net.uk resolved to 212.58.244.71  
ad@haiku:~/src$  
ad@haiku:~/src$ ./GetHostByName-p18 www.nytimes.com  
Name To Be Resolved: www.gtm.nytimes.com  
Name Length in Bytes: 4  
www.gtm.nytimes.com resolved to 170.149.161.130  
ad@haiku:~/src$  
ad@haiku:~/src$ ./GetHostByName-p18 170.149.161.130  
Name To Be Resolved: 170.149.161.130  
Name Length in Bytes: 4  
170.149.161.130 resolved to 170.149.161.130  
ad@haiku:~/src$  
ad@haiku:~/src$
```

# Resolving IP-addresses

```
#include <netdb.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

int main(int argc, char *argv[]) {
    struct hostent* foundhost;
    struct in_addr myaddress;

    /* IPv4 dot-number into binary form (network byte order) */
    inet_aton(argv[1], &myaddress);

    foundhost=gethostbyaddr((const char*)&myaddress, sizeof(myaddress), AF_INET);
    if (foundhost!=NULL){
        printf("IP-address:%s Resolved to: %s\n", argv[1],foundhost->h_name);
        exit(0);
    }
    else {
        printf("IP-address:%s could not be resolved\n",argv[1]);
        exit(1);
    }
}
```

## Resolving IP-addresses

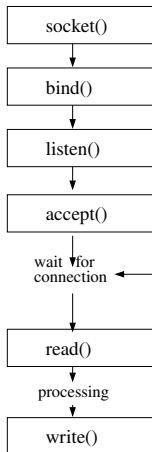
```
ad@haiku:~/src$  
ad@haiku:~/src$ ./GetHostByAddress 128.10.2.166  
IP-address:128.10.2.166 Resolved to: merlin.cs.purdue.edu  
ad@haiku:~/src$  
ad@haiku:~/src$ ./GetHostByAddress 195.134.67.183  
IP-address:195.134.67.183 Resolved to: sydney.di.uoa.gr  
ad@haiku:~/src$
```

- ▶ `gethostbyname()` and `gethostbyaddr()` have been in use.
- ▶ *POSIX.1-2001* suggests instead the use of `getnameinfo()` and `getaddrinfo()` respectively.

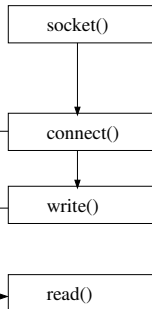
## Our goal

Create the communication endpoint. Use it as a file descriptor.

### Server Process



### Client Process



request for connection establishment

request

response

## Address Format for Sockets

- ▶ An *address* identifies a socket in a **specific communication domain**.
- ▶ Addresses with **different formats** can be passed to the socket functions – all casted to the **generic** sockaddr structure.
- ▶ Internet addresses are defined in <netinet/in.h>.
- ▶ **Specifically in IPv4 Internet domain** (AF\_INET), a socket address is represented by the sockaddr\_in as follows:

```
struct in_addr{
    in_addr_t      s_addr;          /*IPv4  address */
};

struct sockaddr_in{
    sa_family_t     sin_family;     /* address family */
    in_port_t       sin_port;       /* port number    */
    struct in_addr  sin_addr;       /* IPv4 address   */
};
```

- ▶ in\_port\_t data type is uint16\_t (defined in <stdint.h>)
- ▶ in\_addr\_t data type is uint32\_t (defined in <stdint.h>)

# Creating sockets

- ▶ `socket` creates an endpoint for communication
- ▶ returns a descriptor or `-1` on error

```
#include <sys/socket.h>
#include <sys/types.h>
int socket(int domain, int type, int protocol);
```

`domain` communication domain (mostly `AF_INET`)

`type` communication semantics (often `SOCK_STREAM`,  
`SOCK_DGRAM`)

`protocol` Use 0 as typically only one protocol is available

```
if ((sock = socket(AF_INET, SOCK_STREAM, 0)) == -1)
    perror("Socket creation failed!");
```



## Binding sockets to addresses

- ▶ `bind` requests for an address to be assigned to a socket
- ▶ You **must bind** a `SOCK_STREAM` socket to a local address before receiving connections

```
int bind(int socket, const struct sockaddr *address,
         socklen_t address_len);
```

- ▶ We pass a `sockaddr_in` struct as the address that has at least the following members expressed in network byte-order:
  - `sin_family`: address family is `AF_INET` in the Internet domain
  - `sin_addr.s_addr`: address can be a specific IP or `INADDR_ANY`
  - `sin_port`: TCP or UDP port number

# Socket binding example

```
#include <netinet/in.h> /* for sockaddr_in */
#include <sys/socket.h>
#include <sys/types.h>
#include <arpa/inet.h> /* for hton* */

int bind_on_port(int sock, short port) {
    struct sockaddr_in server;
    server.sin_family = AF_INET;
    server.sin_addr.s_addr = htonl(INADDR_ANY);
    server.sin_port = htons(port);
    return bind(sock, (struct sockaddr *) &server, sizeof(server));
}
```

- ▶ INADDR\_ANY is a special address (0.0.0.0) meaning “any address”
- ▶ sock will receive connections from all addresses of the host machine

## listen, accept

```
int listen(int socket, int backlog);
```

- ▶ Listen for connections on a socket
- ▶ At most backlog connections will be queued waiting to be accepted

```
int accept(int socket, struct sockaddr *address,  
           socklen_t *address_len);
```

- ▶ Accepts a connection on a socket
- ▶ Blocks until a client connects/gets-interrupted by a signal
- ▶ Returns new socket descriptor to communicate with client
- ▶ Returns info on clients address through address.  
Pass NULL if you don't care.
- ▶ Value-result address\_len must be set to the amount of space pointed to by address (or NULL).

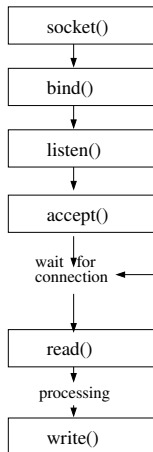
## connect

```
int connect(int socket, struct sockaddr *address,  
            socklen_t address_len);
```

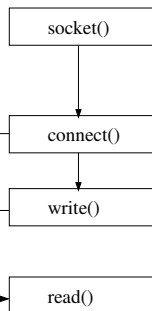
- ▶ When called by a client, a connection is attempted to a listening socket on the server in address. Normally, the server accepts the connection and a communication channel is established.
- ▶ If socket is of type SOCK\_DGRAM, address specifies the peer with which the socket is to be associated (datagrams are sent/received only to/from this peer).

# TCP connection

## Server Process



## Client Process



request for connection establishment

request

response

## Tips and warnings

- ▶ In **Solaris** compile with “-lsocket -lnsl”
- ▶ If a process attempts to write through a socket that has been closed by the other peer, a SIGPIPE signal is received.
- ▶ SIGPIPE is by default fatal, install a signal handler to override this.
- ▶ Use system program netstat to view the status of sockets.

```
ad@linux03:~> netstat -ant
```

- ▶ When a server quits, the listening port remains busy (state TIME\_WAIT) for a while
- ▶ Restarting the server *fails in bind* with “Bind: Address Already in Use”
- ▶ To override this, use setsockopt() to enable SO\_REUSEADDR before you call bind().

TCP server that receives a string and replies with the string capitalized.

```
/*inet_str_server.c: Internet stream sockets server */
#include <stdio.h>
#include <sys/wait.h>          /* sockets */
#include <sys/types.h>         /* sockets */
#include <sys/socket.h>        /* sockets */
#include <netinet/in.h>        /* internet sockets */
#include <netdb.h>             /* gethostbyaddr */
#include <unistd.h>            /* fork */
#include <stdlib.h>            /* exit */
#include <ctype.h>             /* toupper */
#include <signal.h>            /* signal */
void child_server(int newsock);
void perror_exit(char *message);
void sigchld_handler (int sig);

void main(int argc, char *argv[]) {
    int                port, sock, newsock;
    struct sockaddr_in server, client;
    socklen_t clientlen;
    struct sockaddr *serverptr=(struct sockaddr *)&server;
    struct sockaddr *clientptr=(struct sockaddr *)&client;
```

```

struct hostent *rem;
if (argc != 2) {
    printf("Please give port number\n");exit(1);}
port = atoi(argv[1]);
/* Reap dead children asynchronously */
signal(SIGCHLD, sigchld_handler);
/* Create socket */
if ((sock = socket(AF_INET, SOCK_STREAM, 0)) < 0)
    perror_exit("socket");
server.sin_family = AF_INET;          /* Internet domain */
server.sin_addr.s_addr = htonl(INADDR_ANY);
server.sin_port = htons(port);        /* The given port */
/* Bind socket to address */
if (bind(sock, serverptr, sizeof(server)) < 0)
    perror_exit("bind");
/* Listen for connections */
if (listen(sock, 5) < 0) perror_exit("listen");

```

ΣΥΧΝΑ ΠΡΙΝ ΤΟ bind

```

int reuse = 1;
if (setsockopt(sock, SOL_SOCKET, SO_REUSEADDR,
                (const char*)&reuse, sizeof(reuse)) < 0)
    perror("setsockopt(SO_REUSEADDR) failed");

```



```

printf("Listening for connections to port %d\n", port);
while (1) {
    /* accept connection */
    if ((newsock = accept(sock, clientptr, &clientlen)) < 0) perror_exit("
        accept");
    /* Find client's address */
    if ((rem = gethostbyaddr((char *) &client.sin_addr.s_addr, sizeof(client
        .sin_addr.s_addr), client.sin_family)) == NULL) {
        perror("gethostbyaddr"); exit(1);}
    printf("Accepted connection from %s\n", rem->h_name);
    // printf("Accepted connection\n");
    switch (fork()) {    /* Create child for serving client */
    case -1:             /* Error */
        perror("fork"); break;
    case 0:             /* Child process */
        close(sock); child_server(newsock);
        exit(0);
    }
    close(newsock); /* parent closes socket to client */
}

```

```

    }
}

void child_server(int newsock) {
    char buf[1];
    while(read(newsock, buf, 1) > 0) { /* Receive 1 char */
        putchar(buf[0]); /* Print received char */
        /* Capitalize character */
        buf[0] = toupper(buf[0]);
        /* Reply */
        if (write(newsock, buf, 1) < 0)
            perror_exit("write");
    }
    printf("Closing connection.\n");
    close(newsock); /* Close socket */
}

/* Wait for all dead child processes */
void sigchld_handler (int sig) {
    while (waitpid(-1, NULL, WNOHANG) > 0);
}

void perror_exit(char *message) {
    perror(message);
    exit(EXIT_FAILURE);
}

```

## TCP client example. (definitions)

```
/* inet_str_client.c: Internet stream sockets client */
#include <stdio.h>
#include <sys/types.h>          /* sockets */
#include <sys/socket.h>         /* sockets */
#include <netinet/in.h>         /* internet sockets */
#include <unistd.h>             /* read, write, close */
#include <netdb.h>              /* gethostbyaddr */
#include <stdlib.h>             /* exit */
#include <string.h>             /* strlen */

void perror_exit(char *message);

void main(int argc, char *argv[]) {
    int          port, sock, i;
    char         buf[256];
    struct sockaddr_in server;
    struct sockaddr *serverptr = (struct sockaddr*)&server;
    struct hostent *rem;
    if (argc != 3) {
        printf("Please give host name and port number\n");
        exit(1);}
}
```

## TCP client example. (connection)

```
/* Create socket */
if ((sock = socket(AF_INET, SOCK_STREAM, 0)) < 0)
    perror_exit("socket");
/* Find server address */
if ((rem = gethostbyname(argv[1])) == NULL) {
    perror("gethostbyname"); exit(1);
}
port = atoi(argv[2]); /*Convert port number to integer*/
server.sin_family = AF_INET; /* Internet domain */
memcpy(&server.sin_addr, rem->h_addr, rem->h_length);
server.sin_port = htons(port); /* Server port */
/* Initiate connection */
if (connect(sock, serverptr, sizeof(server)) < 0)
    perror_exit("connect");
printf("Connecting to %s port %d\n", argv[1], port);
```

## TCP client example. (transfer loop)

```
do {
    printf("Give input string: ");
    fgets(buf, sizeof(buf), stdin); /* Read from stdin */
    for(i=0; buf[i] != '\0'; i++) { /* For every char */
        /* Send i-th character */
        if (write(sock, buf + i, 1) < 0)
            perror_exit("write");
        /* receive i-th character transformed */
        if (read(sock, buf + i, 1) < 0)
            perror_exit("read");
    }
    printf("Received string: %s", buf);
} while (strcmp(buf, "END\n") != 0); /* Finish on "end" */
close(sock); /* Close socket and exit */
}

void perror_exit(char *message)
{
    perror(message);
    exit(EXIT_FAILURE);
}
```

# Execution

Server on linux02:

```
ad@linux02:~> ./server 9002
Listening for connections to port 9002
Accepted connection from linux03.di.uoa.gr
Hello world
EnD
Closing connection.
```

Client on linux03:

```
ad@linux03:~> ./client linux02.di.uoa.gr 9002
Connecting to linux02.di.uoa.gr port 9002
Give input string: Hello world
Received string: HELLO WORLD
Give input string: EnD
Received string: END
ad@linux03:~>
```

## More useful functions

`shutdown` shut down part of a full-duplex connection

```
int shutdown(int socket, int how);
```

Can be used to tell server that we have sent the whole request.

`getsockname` returns the current address to which the socket is bound with using the buffer pointed to by address.

```
int getsockname(int socket,
                struct sockaddr *address,
                socklen_t *address_len);
```

`getpeername` get the name (address) of the peer connected to a socket; useful if a server has called a fork/exec combination and only the socket is known.

```
int getpeername(int socket,
                struct sockaddr *address,
                socklen_t *address_len);
```

# Parsing and Printing Addresses

`inet_ntoa` Convert struct `in_addr` to printable form 'a.b.c.d'

`inet_addr` Convert IP address string in '.' notation to 32bit network address

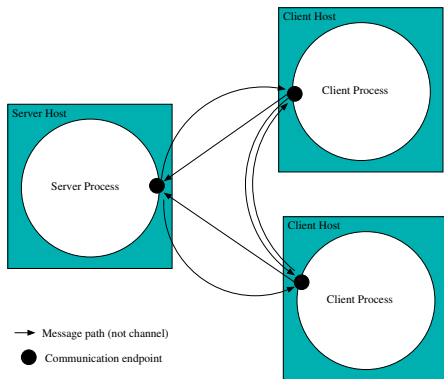
`inet_ntop` Convert address from network format to printable presentation format

`inet_pton` Convert presentation format address to network format

**Bonus:** `inet_ntop` and `inet_pton` also work with IPv6!

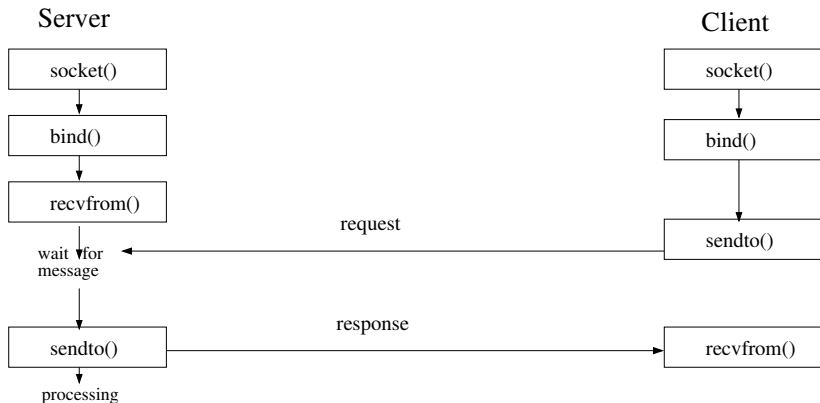


# Internet User Datagram Protocol (UDP)



- ▶ No Connections: Think postcards, not telephone.
- ▶ Datagrams (messages) exchanged.
- ▶ Datagrams *either* arrive (possibly out of order) *or* get lost!

# UDP communication



## sendto, recvfrom

```
ssize_t sendto(int sock, void *buff, size_t length,  
               int flags, struct sockaddr *dest_addr,  
               socklen_t dest_len);
```

- ▶ Send a message to a socket
- ▶ Similar to write() & send() but designates destination

```
ssize_t recvfrom(int socket, void *buff, size_t length,  
                 int flags, struct sockaddr *address,  
                 socklen_t *address_len);
```

- ▶ Receive a message from a socket
- ▶ Similar to read() & recv() but designates source address
- ▶ address\_len is value-result and must be initialized to the size of the buffer pointed to by the address pointer
- ▶ last two arguments can be NULL

Usually flags = 0 ; rarely used (ex. out of band data)

## A simple echoing UDP server

Client on linux03 (along with the input to send out):

```
ad@linux03:~> fortune | ./inet_dgm_client linux02 59579
Hlade's Law:
    If you have a difficult task, give it to a lazy person --
    they will find an easier way to do it.
ad@linux03:~>
```

Server on linux02 (along with the bytes to receive from client):

```
ad@linux02:~> ./inet_dgm_server
Socket port: 59579
Received from linux03: Hlade's Law:
Received from linux03:  If you have a difficult task, give it to a lazy person
--
Received from linux03:  they will find an easier way to do it.
```

```

/* inet_dgr_server.c: Internet datagram sockets server */
#include <sys/types.h> /* sockets */
#include <sys/socket.h> /* sockets */
#include <netinet/in.h> /* Internet sockets */
#include <netdb.h> /* gethostbyaddr */
#include <arpa/inet.h> /* inet_ntoa */
#include <stdio.h>
#include <stdlib.h>
void perror_exit(char *message);

char *name_from_address(struct in_addr addr) {
    struct hostent *rem; int asize = sizeof(addr.s_addr);
    if((rem = gethostbyaddr(&addr.s_addr, asize, AF_INET)))
        return rem->h_name; /* reverse lookup success */
    return inet_ntoa(addr); /* fallback to a.b.c.d form */
}

void main() {
    int n, sock; unsigned int serverlen, clientlen;
    char buf[256], *clientname;
    struct sockaddr_in server, client;
    struct sockaddr *serverptr = (struct sockaddr*) &server;
    struct sockaddr *clientptr = (struct sockaddr*) &client;
    /* Create datagram socket */

```

```

if ((sock = socket(AF_INET, SOCK_DGRAM, 0)) < 0)
    perror_exit("socket");
/* Bind socket to address */
server.sin_family = AF_INET;           /* Internet domain */
server.sin_addr.s_addr = htonl(INADDR_ANY);
server.sin_port = htons(0);           /* Autoselect a port */
serverlen = sizeof(server);
if (bind(sock, serverptr, serverlen) < 0)
    perror_exit("bind");
/* Discover selected port */
if (getsockname(sock, serverptr, &serverlen) < 0)
    perror_exit("getsockname");
printf("Socket port: %d\n", ntohs(server.sin_port));
while(1) { clientlen = sizeof(client);
    /* Receive message */
    if ((n = recvfrom(sock, buf, sizeof(buf), 0, clientptr, &clientlen)) <
        0)
        perror("recvfrom");
    buf[sizeof(buf)-1]='\0'; /* force str termination */
    /* Try to discover client's name */
    clientname = name_from_address(client.sin_addr);
    printf("Received from %s: %s\n", clientname, buf);
    /* Send message */
    if (sendto(sock, buf, n, 0, clientptr, clientlen)<0)
        perror_exit("sendto");
}}

void perror_exit(char *message)
{
    perror(message);
    exit(EXIT_FAILURE);
}

```

```

/* inet_dgr_client.c: Internet datagram sockets client */
#include <sys/types.h> /* sockets */
#include <sys/socket.h> /* sockets */
#include <netinet/in.h> /* Internet sockets */
#include <netdb.h> /* gethostbyname */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

void main(int argc, char *argv[]) {
    int sock; char buf[256]; struct hostent *rem;
    struct sockaddr_in server, client;
    unsigned int serverlen = sizeof(server);
    struct sockaddr *serverptr = (struct sockaddr *) &server;
    struct sockaddr *clientptr = (struct sockaddr *) &client;
    if (argc < 3) {
        printf("Please give host name and port\n"); exit(1);}
    /* Create socket */
    if ((sock = socket(AF_INET, SOCK_DGRAM, 0)) < 0) {
        perror("socket"); exit(1); }
    /* Find server's IP address */
    if ((rem = gethostbyname(argv[1])) == NULL) {
        perror("gethostbyname"); exit(1); }
}

```

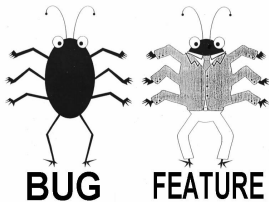
```

/* Setup server's IP address and port */
server.sin_family = AF_INET;           /* Internet domain */
memcpy(&server.sin_addr, rem->h_addr, rem->h_length);
server.sin_port = htons(atoi(argv[2]));
/* Setup my address */
client.sin_family = AF_INET;           /* Internet domain */
client.sin_addr.s_addr=htonl(INADDR_ANY); /*Any address*/
client.sin_port = htons(0);            /* Autoselect port */
/* Bind my socket to my address*/
if (bind(sock, clientptr, sizeof(client)) < 0) {
    perror("bind"); exit(1); }
/* Read continuously messages from stdin */
while (fgets(buf, sizeof buf, stdin)) {
    buf[strlen(buf)-1] = '\0';          /* Remove '\n' */
    if (sendto(sock, buf, strlen(buf)+1, 0, serverptr, serverlen) < 0) {
        perror("sendto"); exit(1); }   /* Send message */
    bzero(buf, sizeof buf);            /* Erase buffer */
    if (recvfrom(sock, buf, sizeof(buf), 0, NULL, NULL) < 0) {
        perror("recvfrom"); exit(1); }  /* Receive message */
    printf("%s\n", buf);
}
}

```



- ▶ Everything looks good and runs ok BUT there is a **BUG!**
- ▶ Remember that UDP is *unreliable*



## rlsd: a remote file/dir-listing server

Server on linux02:

```
ad@linux02:~> ./rlsd
```

Client on linux03:

```
ad@linux03:~> ./rls linux02.di.uoa.gr /usr/share/dict
README
connectives
propernames
web2
web2a
words
ad@linux03:~>
```

`rlsd.c` remote ls server (TCP)

`fdopen` allows buffered I/O by opening socket as file stream

```
/* rlsd.c - a remote ls server - with paranoia */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>
#include <time.h>
#include <string.h>
#include <ctype.h>
#define PORTNUM 15000 /* rlsd listens on this port */

void perror_exit(char *msg);
void sanitize(char *str);
```

```

int main(int argc, char *argv[]) {
    struct sockaddr_in myaddr; /* build our address here */
    int c, lsock, csock; /* listening and client sockets */
    FILE *sock_fp; /* stream for socket IO */
    FILE *pipe_fp; /* use popen to run ls */
    char dirname[BUFSIZ]; /* from client */
    char command[BUFSIZ]; /* for popen() */

    /** create a TCP a socket **/
    if ((lsock = socket( AF_INET, SOCK_STREAM, 0)) < 0)
        perror_exit( "socket" );
    /** bind address to socket. **/
    myaddr.sin_addr.s_addr = htonl(INADDR_ANY);
    myaddr.sin_port = htons(PORTNUM); /*port to bind socket*/
    myaddr.sin_family = AF_INET; /* internet addr family */
    if(bind(lsock,(struct sockaddr *)&myaddr, sizeof(myaddr)))
        perror_exit( "bind" );
    /** listen for connections with Qsize=5 **/
    if ( listen(lsock, 5) != 0 )
        perror_exit( "listen" );

```

```

while ( 1 ){ /* main loop: accept - read - write */
    /* accept connection, ignore client address */
    if ( (csock = accept(lsock, NULL, NULL)) < 0 )
        perror_exit("accept");
    /* open socket as buffered stream */
    if ((sock_fp = fdopen(csock,"r+")) == NULL)
        perror_exit("fdopen");
    /* read dirname and build ls command line */
    if (fgets(dirname, BUFSIZ, sock_fp) == NULL)
        perror_exit("reading dirname");
    sanitize(dirname); /* clear wild characters */
    snprintf(command, BUFSIZ, "ls %s", dirname);
    /* Invoke ls through popen */
    if ((pipe_fp = popen(command, "r")) == NULL )
        perror_exit("popen");
    /* transfer data from ls to socket */
    while( (c = getc(pipe_fp)) != EOF )
        putc(c, sock_fp);
    pclose(pipe_fp);
    fclose(sock_fp);
}
return 0;
}

```

```
/* it would be very bad if someone passed us an dirname like  
 * "; rm *" and we naively created a command "ls ; rm *".  
 * So..we remove everything but slashes and alphanumerics.  
 */  
void sanitize(char *str)  
{  
    char *src, *dest;  
    for ( src = dest = str ; *src ; src++ )  
        if ( *src == '/' || isalnum(*src) )  
            *dest++ = *src;  
    *dest = '\0';  
}  
  
/* Print error message and exit */  
void perror_exit(char *message)  
{  
    perror(message);  
    exit(EXIT_FAILURE);  
}
```

`rls.c` sends a directory name to `rlsd` and reads back a directory listing (TCP)

`write_all` guarantees to send all the bytes requested, provided no error occurs, by repeatedly calling `write()`

```
#include <sys/types.h> /* sockets */
#include <sys/socket.h> /* sockets */
#include <netinet/in.h> /* internet sockets */
#include <netdb.h> /* gethostbyname */
#define PORTNUM 15000
#define BUFFSIZE 256
void perror_exit(char *msg);

/* Write() repeatedly until 'size' bytes are written */
int write_all(int fd, void *buff, size_t size) {
    int sent, n;
    for(sent = 0; sent < size; sent+=n) {
        if ((n = write(fd, buff+sent, size-sent)) == -1)
            return -1; /* error */
    }
    return sent;
}
```

```

int main(int argc, char *argv[]) {
    struct sockaddr_in servadd; /* The address of server */
    struct hostent *hp;         /* to resolve server ip */
    int sock, n_read;           /* socket and message length */
    char buffer[BUFSIZE];       /* to receive message */

    if ( argc != 3 ) {
        puts("Usage: rls <hostname> <directory>"); exit(1);}
    /* Step 1: Get a socket */
    if ((sock = socket(AF_INET, SOCK_STREAM, 0)) == -1 )
        perror_exit( "socket" );
    /* Step 2: lookup server's address and connect there */
    if ((hp = gethostbyname(argv[1])) == NULL) {
        perror("gethostbyname"); exit(1);}
    memcpy(&servadd.sin_addr, hp->h_addr, hp->h_length);
    servadd.sin_port = htons(PORTNUM); /* set port number */
    servadd.sin_family = AF_INET ;     /* set socket type */
    if (connect(sock, (struct sockaddr*) &servadd, sizeof(servadd)) !=0)
        perror_exit( "connect" );
    /* Step 3: send directory name + newline */

```



```
if ( write_all(sock, argv[2], strlen(argv[2])) == -1)
    perror_exit("write");
if ( write_all(sock, "\n", 1) == -1 )
    perror_exit("write");
/* Step 4: read back results and send them to stdout */
while( (n_read = read(sock, buffer, BUFSIZE)) > 0 )
    if (write_all(STDOUT_FILENO, buffer, n_read)<n_read)
        perror_exit("fwrite");
close(sock);
return 0;
}
```

# Select

When listening for connections on `sockfd`, we can probe multiple connections for input:

- ▶ Use function `select()`
- ▶ Select probes a bitmap (`fd_set`) of active file descriptors for input
- ▶ Input to `sockfd` means that there is a new connection attempt => accept the connection
- ▶ Input from another `fd` means that data can be read from it
- ▶ When `read()` from an `fd` returns 0, we have EOF => close the corresponding connection
- ▶ Useful for multithreaded server model

# Select

```
#include <stdio.h>
#include <errno.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>

#define PORT    5555
#define MAXMSG  512

int
read_from_client (int filedes)
{
    char buffer[MAXMSG];
    int nbytes;

    nbytes = read (filedes, buffer, MAXMSG);
    if (nbytes < 0)
    {
        /* Read error. */
        perror ("read");
        exit (EXIT_FAILURE);
    }
    else if (nbytes == 0)
        /* End-of-file. */
        return -1;
    else
    {
        /* Data read. */
        fprintf (stderr, "Server: got message: `'%s'\n", buffer);
        return 0;
    }
}
```

## Select

```
int
main (void)
{
    extern int make_socket (uint16_t port);
    int sock;
    fd_set active_fd_set, read_fd_set;
    int i;
    struct sockaddr_in clientname;
    size_t size;

    /* Create the socket and set it up to accept connections. */
    sock = make_socket (PORT);
    if (listen (sock, 1) < 0)
    {
        perror ("listen");
        exit (EXIT_FAILURE);
    }
}
```

## Select

```
/* Initialize the set of active sockets. */
FD_ZERO (&active_fd_set);
FD_SET (sock, &active_fd_set);

while (1)
{
    /* Block until input arrives on one or more active sockets. */
    read_fd_set = active_fd_set;
    if (select (FD_SETSIZE, &read_fd_set, NULL, NULL, NULL) < 0)
    {
        perror ("select");
        exit (EXIT_FAILURE);
    }

    /* Service all the sockets with input pending. */
    for (i = 0; i < FD_SETSIZE; ++i)
        if (FD_ISSET (i, &read_fd_set))
        {
            . . .
        }
    . . .
}
```

# Select

```
/* Service all the sockets with input pending. */
for (i = 0; i < FD_SETSIZE; ++i)
    if (FD_ISSET (i, &read_fd_set))
    {
        if (i == sock)
        {
            /* Connection request on original socket. */
            int new;
            size = sizeof (clientname);
            new = accept (sock,
                          (struct sockaddr *) &clientname,
                          &size);

            if (new < 0)
            {
                perror ("accept");
                exit (EXIT_FAILURE);
            }

            fprintf (stderr,
                    "Server: connect from host %s, port %hd.\n",
                    inet_ntoa (clientname.sin_addr),
                    ntohs (clientname.sin_port));
            FD_SET (new, &active_fd_set);
        }
    }
```

## Select

```
    }  
else  
{  
    /* Data arriving on an already-connected socket. */  
    if (read_from_client (i) < 0)  
    {  
        close (i);  
        FD_CLR (i, &active_fd_set);  
    }  
    }  
}  
}  
}
```